


Linux Security Modules

Joseph Hallett



How do we get Linux to Grade B?

TSEC says Linux just offers discretionary access control

- Grade C (pushing into B with things like AppArmour)

How do we bring MAC policies to Linux?

- Needs to be small and testable and not impact people who don't want to use a MAC policy on their systems...

(This is/was a real research problem)

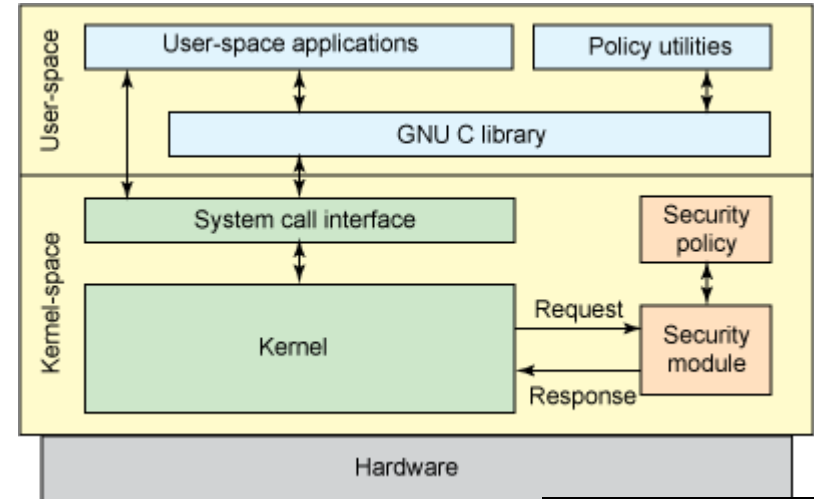


LSM Framework

Linux Security Module framework
implements a reference monitor for Linux

Dynamically loadable kernel module hooks
into system call checks

Framework is verified, modules are (in
theory) small and verifiable




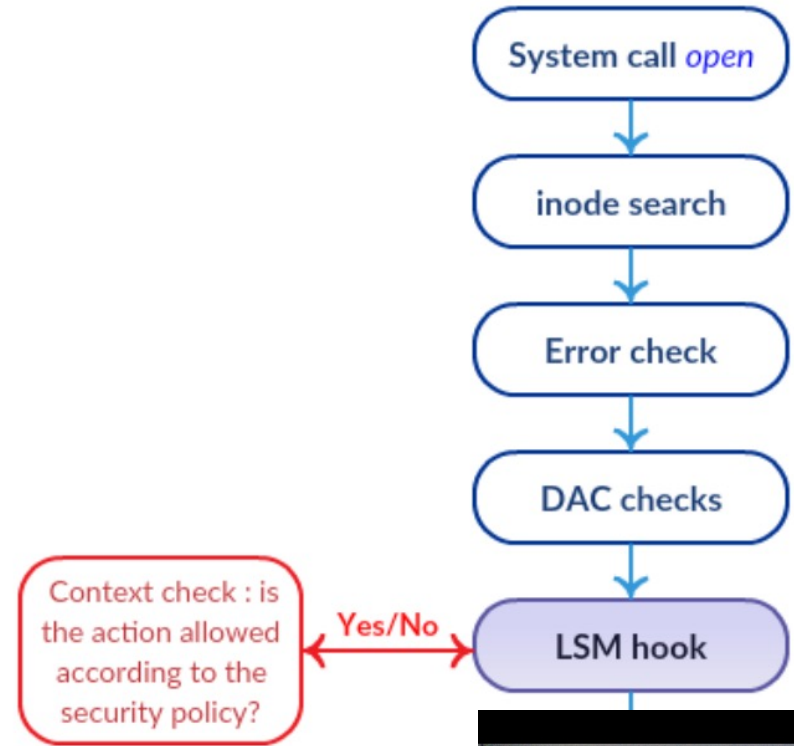
LSM Framework

Security hooks are invoked on the path between any subject action and an object

- For example, *processes* and *inodes*

Hook function returns access decision:

- 0: access granted
- ENOMEM: no memory available
- EPERM: not enough privileges




SELinux

LSM framework developed (initially) by the NSA

- Open sourced in 2000
- Merged into Linux Kernel 2.6.0 in 2003
- Available in Fedora (Core 2) since 2004

Based around type-based enforcement and RBAC (role based access controls)




SELinux Hook Example

security/selinux/hooks.c

bristol.ac.uk

```
3680  /* file security operations */
3681
3682  static int selinux_revalidate_file_permission(struct file *file, int mask)
3683  {
3684      const struct cred *cred = current_cred();
3685      struct inode *inode = file_inode(file);
3686
3687      /* file_mask_to_av won't add FILE__WRITE if MAY_APPEND is set */
3688      if ((file->f_flags & O_APPEND) && (mask & MAY_WRITE))
3689          mask |= MAY_APPEND;
3690
3691      return file_has_perm(cred, file,
3692                          file_mask_to_av(inode->i_mode, mask));
3693  }
3694
3695  static int selinux_file_permission(struct file *file, int mask)
3696  {
3697      struct inode *inode = file_inode(file);
3698      struct file_security_struct *fsec = selinux_file(file);
3699      struct inode_security_struct *isec;
3700      u32 sid = current_sid();
3701
3702      if (!mask)
3703          /* No permission to check. Existence test. */
3704          return 0;
3705
3706      isec = inode_security(inode);
3707      if (sid == fsec->sid && fsec->isid == isec->sid)
3708          fsec->pseqno == avc_policy_seqno(&selinux_policy);
3709      /* No change since file_open check. */
3710      return 0;
3711
3712      return selinux_revalidate_file_permission(file, mask);
3713  }
```



SELinux Hook Example

security/selinux/hooks.c

bristol.ac.uk

```
1759 /* Check whether a task can use an open file descriptor to
1760    access an inode in a given way. Check access to the
1761    descriptor itself, and then use dentry_has_perm to
1762    check a particular permission to the file.
1763    Access to the descriptor is implicitly granted if it
1764    has the same SID as the process. If av is zero, then
1765    access to the file is not checked, e.g. for cases
1766    where only the descriptor is affected like seek. */
1767 static int file_has_perm(const struct cred *cred,
1768                          struct file *file,
1769                          u32 av)
1770 {
1771     struct file_security_struct *fsec = selinux_file(file);
1772     struct inode *inode = file_inode(file);
1773     struct common_audit_data ad;
1774     u32 sid = cred_sid(cred);
1775     int rc;
1776
1777     ad.type = LSM_AUDIT_DATA_FILE;
1778     ad.u.file = file;
1779
1780     if (sid != fsec->sid) {
1781         rc = avc_has_perm(&selinux_state,
1782                           sid, fsec->sid,
1783                           SECCLASS_FD,
1784                           FD_USE,
1785                           &ad);
1786
1787         if (rc)
1788             goto out;
1789     }
1790
1791 #ifdef CONFIG_BPF_SYSCALL
1792     rc = bpf_fd_pass(file, cred_sid(cred));
1793     if (rc)
1794         return rc;
1795 #endif
1796
1797     /* av is zero if only checking
1798     rc = 0;
1799     if (av)
1800         rc = inode_has_perm(cred
1801 out:
1802     return rc;
1803 }
```



SELinux Hooks

Management hooks

- Called to handle object lifecycle
- E.g. `security_inode_allocate` or `security_inode_free`
- Used to manage security information

Path-based hooks

- Related to pathnames (used to implement TSEC B1)

Object-based hooks

- Path kernel structure corresponding to objects (used to implement TSEC B2)



SELinux Pseudo Filesystem


Need a mechanism to interact with SELinux from userland

- To load policies and configuration
- To get audit data

Standard mountable virtual filesystem (like /proc)

- security/selinux/selinuxfs.c

```
124 | #define TMPBUFLen      12
125 | static ssize_t sel_read_enforce(struct file *filp, char __user *buf,
126 |                               size_t count, loff_t *ppos)
127 | {
128 |     struct selinux_fs_info *fsi = file_inode(filp)->i_sb->s_fs_info;
129 |     char tmpbuf[TMPBUFLen];
130 |     ssize_t length;
131 |
132 |     length = scnprintf(tmpbuf, TMPBUFLen, "%d",
133 |                       enforcing_enabled(fsi->state)
134 |     return simple_read_from_buffer(buf, count, ppos,
135 | }
```



SELinux Policies

All subjects get labeled with a security context

- User
- Domain
- Role

Rules describe what each *subject* domain can do with an *object* domain...

They get a bit complicated...



SELinux Policy Example

`/etc/passwd`

User information readable by any user

`/etc/shadow`

Password information readable only by root(ish)



SELinux Policy Example

```
/* Normal users are allowed to read normal files */  
allow user_t public_t : file read
```



SELinux Policy Example

```
/* Normal users are allowed to read normal files */  
allow user_t public_t : file read
```

```
/* Users in the password_t domain can r/w files in the  
   password_data_t domain */  
allow passwd_t passwd_data_t : file {read write}
```

```
/* Allow users to actually run the password program, and transition their domain */  
allow user_t passwd_exec_t : file execute  
allow user_t passwd_t : process transition  
type transition user_t passwd_exec_t : process passwd_t
```



It's kinda complex!

5 lines of policy just to control access to 2 files

- Kinda makes sense...?

NSA gives a reference policy

- Hard to live with, and long (~20,000 rules)
- `echo 0 >/sys/fs/selinux/enforce`

Android implements its permissions system using its own SELinux policy

- Also long and confusing...

Rule design is really hard



SELinux Coloring Book

Rule design is hard...

So lets try and make it easy!

https://people.redhat.com/duffy/selinux/selinux-coloring-book_A4-Stapled.pdf

bristol.ac.uk

